

# Term weighting and Vector space retrieval

# Vector Space Model

---

- It represents documents and queries as vectors of features representing terms
- features are assigned some numerical value that is usually some function of frequency of terms
- Ranking algorithm compute similarity between document and query vectors to yield a retrieval score to each document.

## Documents as vectors

---

- Each doc  $d$  is viewed as a vector of  $tf \times idf$  values, one component for each term
- So we have a vector space
  - terms are axes
  - docs live in this space

# Vector Space Model

---

- Given a finite set of  $n$  documents:

$$D = \{d_1, d_2, \dots, d_j, \dots, d_n\}$$

and a finite set of  $m$  terms:

$$T = \{t_1, t_2, \dots, t_i, \dots, t_m\}$$

- Each document will be represented by a column vector of weights as follows:

$$(w_{1j}, w_{2j}, w_{3j}, \dots, w_{ij}, \dots, w_{mj})^t$$

$w_{ij}$  is the weight of term  $t_i$  in document  $d_j$ .

# Vector Space Model

---

The document collection as a whole will be represented by an  $m \times n$  term-document matrix as:

$$\begin{pmatrix} W_{11} & W_{12} & \dots & W_{1j} & \dots & W_{1n} \\ W_{21} & W_{22} & \dots & W_{2j} & \dots & W_{2n} \\ W_{i1} & W_{i2} & \dots & W_{ij} & \dots & W_{in} \\ W_{m1} & W_{m2} & \dots & W_{mj} & \dots & W_{mn} \end{pmatrix}$$

## Example: Vector Space Model

---

D1 = Information retrieval is concerned with the organization, storage, retrieval and evaluation of information relevant to user's query.

D2 = A user having an information need formulates a request in the form of query written in natural language.

D3 = The retrieval system responds by retrieving document that seems relevant to the query.

## Example: Vector Space Model

---

Let the weights be assigned based on the frequency of the term within the document.

The term – document matrix is:

$$\begin{pmatrix} 2 & 1 & 0 \\ 2 & 0 & 1 \\ 1 & 1 & 1 \end{pmatrix}$$

## Vector Space Model

---

- Raw term frequency approach gives too much importance to the absolute values of various coordinates of each document



---

Consider two document vectors

$$(2, 2, 1)^t$$

$$(4, 4, 2)^t$$

The documents look similar except the differences in magnitude of term weights.

## Normalizing term weights

---

- To reduce the importance of the length of document vectors we normalize document vectors
- Normalization changes all the vectors to a standard length.

We can convert document vectors to unit length by dividing each dimension by the overall length of the vector.

---

Normalizing the term-document matrix:

$$\begin{pmatrix} 2 & 1 & 0 \\ 2 & 0 & 1 \\ 1 & 1 & 1 \end{pmatrix}$$

We get

$$\begin{pmatrix} 0.67 & 0.71 & 0 \\ 0.67 & 0 & 0.71 \\ 0.33 & 0.71 & 0.71 \end{pmatrix}$$

Elements of each column are divided by the length of the column vector ( $\sqrt{\sum_i w_{ij}^2}$ )

# Term weighting

---

## Postulates

1. The more a document contains a given word the more that document is about a concept represented by that word.
2. The less a term occurs in particular document in a collection, the more discriminating that term is.

## Term weighting

---

- The first factor simply means that terms that occur more frequently represent its meaning more strongly than those occurring less frequently
- The second factor considers term distribution across the document collection.

## Term weighting

---

→ a measure that favors terms appearing in fewer documents is required

The fraction  $n/n_i$ , exactly gives this measure where,

$n$  is the total number of the document in the collection

&  $n_i$  is the number of the document in which term  $i$  occurs

# Term weighting

---

- As the number of documents in any collection is usually large, log of this measure is usually taken, resulting in the following form of inverse document frequency (idf) term weight:

$$\text{idf}_i = \log \left( \frac{n}{n_i} \right)$$

## Tf-idf weighting scheme

---

$$w_{ij} = tf_{ij} \times \log\left(\frac{n}{n_i}\right)$$

tf - document specific statistic

idf - is global statistic and attempts to include distribution of term across document collection.



## Tf-idf weighting scheme

---

- The term frequency (tf) component is document specific statistic that measures the importance of term within the document
- The inverse document frequency (idf) is global statistic and attempts to include distribution of term across document collection.

## Tf-idf weighting scheme

---

Example: Computing tf-idf weight( total docs=100)

term	frequency (tf)	Document frequency ( $n_j$ )	idf ( $\log(n/n_j)$ )	Weight (tf x idf)
Tornado	4	15	0.824	3.296
Swirl	1	20	0.699	0.699
Wind	1	40	0.398	0.389

## Normalizing tf and idf factors

---

- by dividing the term frequency by the frequency of the most frequent term in the document
- idf can be normalized by dividing it by the logarithm of the collection size (n).

$$w_{ij} = \frac{tf_{ij}}{\max(tf_{ij})} \times \log\left(\frac{n}{n_i}\right) / \log(n)$$

# Term weighting schemes

---

- A third factor that may affect weighting function is the document length
- the weighting schemes can thus be characterized by the following three factors:
  1. Within-document frequency or term frequency
  2. Collection frequency or inverse document frequency
  3. Document length

- 
- term weighting scheme can be represented by a triple ABC
    - A - tf component
    - B - idf component
    - & C - length normalization component.

- 
- Different combinations of options can be used to represent document and query vectors.
  - The retrieval model themselves can be represented by a pair of triples like `nnn.nnn` (doc = “nnn”, query = “nnn”)

# Options for the three weighting factors

---

- **Term frequency within document (A)**

n Raw term frequency  $tf = tf_{ij}$

b  $tf = 0$  or  $1$  (binary weight)

a  $tf = 0.5 + 0.5 \left( \frac{tf_{ij}}{\max tf \text{ in } D_j} \right)$  Augmented term frequency

l  $tf = \ln(tf_{ij}) + 1.0$  Logarithmic term frequency

- **Inverse Document frequency (B)**

n  $wt = tf$  no conversion

t Multiply  $tf$  with  $idf$

# Options for the three weighting factors

---

- **Document length (C)**

n             $w_{ij} = wt$             (no conversion)

c             $w_{ij}$  is obtained by dividing each  $wt$   
by  $\text{sqrt}(\text{sum of}(wts \text{ squared}))$



# Indexing Algorithm

---

Step 1. Tokenization: This extracts individual terms (words) in the document, converts all the words in the lower case and removes punctuation marks. The output of the first stage is a representation of the document as a stream of terms.

Step 2. Stop word elimination: Removes words that appear more frequently in the document collection.

Step 3. Stemming: reduce remaining terms to their linguistic root, to get index terms.

Step 4. Term weighting: Assigns weights to term according to their importance in the document, in the collection or some combination of both.

# Example: Document Representation

---

Document 1: Vector space model

Document 2: Probabilistic retrieval model

Document 3: Intelligent techniques in information retrieval

Stemmed terms	Document 1	Document 2	Document 3
inform	0	0	1
intellig	0	0	1
model	1	1	0
probabilist	0	1	0
retriev	0	1	1
space	1	0	0
technique	0	0	1
vector	1	0	0


## Why turn docs into vectors?

---

- First application: Query-by-example
  - Given a doc  $d$ , find others “like” it.
- Now that  $d$  is a vector, find vectors (docs) “near” it.

# Intuition

---



Postulate: Documents that are “close together” in the vector space talk about the same things.

## Desiderata for proximity

---

- If  $d_1$  is near  $d_2$ , then  $d_2$  is near  $d_1$ .
- If  $d_1$  near  $d_2$ , and  $d_2$  near  $d_3$ , then  $d_1$  is not far from  $d_3$ .
- No doc is closer to  $d$  than  $d$  itself.

## First cut

---


- Idea: Distance between  $d_1$  and  $d_2$  is the length of the vector  $|d_1 - d_2|$ .
  - Euclidean distance
- Why is this not a great idea?
  - Short documents would be more similar to each other by virtue of length, not topic
- However, we can implicitly normalize by looking at *angles* instead



## Cosine similarity

- Distance between vectors  $d_1$  and  $d_2$  captured by the cosine of the angle  $\theta$  between them.
- Note – this is *similarity*, not distance

$$\text{sim}(d_j, d_k) = \frac{\sum_{i=1}^n w_{i,j} w_{i,k}}{\sqrt{\sum_{i=1}^n w_{i,j}^2} \sqrt{\sum_{i=1}^n w_{i,k}^2}}$$



## Cosine similarity

---

- A vector can be *normalized* (given a length of 1) by dividing each of its components by its length

$$\|\mathbf{x}\|_2 = \sqrt{\sum_i x_i^2}$$

- This maps vectors onto the unit sphere:
- Then,
- Longer documents don't get more weight



# Cosine similarity

---

- Cosine of angle between two vectors
- The denominator involves the lengths of the vectors.



Normalization

## Normalized vectors

---

- For normalized vectors, the cosine is simply the dot product:

# Queries in the vector space model

---


## Central idea: the query as a vector:

- We regard the query as short document
- We return the documents ranked by the closeness of their vectors to the query, also represented as a vector.

- Note that  $d$  is very general

$$\text{sim}(d_j, d_q) = \frac{\sum_{i=1}^n w_{i,j} w_{i,q}}{\sqrt{\sum_{i=1}^n w_{i,j}^2} \sqrt{\sum_{i=1}^n w_{i,q}^2}}$$

# Similarity Measures



- cosine similarity  $sim(d_j, q_k) = \frac{(d_j, q_k)}{\|d_j\| \|q_k\|} = \frac{\sum_{i=1}^m w_{ij} \times w_{ik}}{\sqrt{\sum_{i=1}^m w_{ik}^2} \times \sqrt{\sum_{i=1}^m w_{ij}^2}}$

---

Let

$$D = (0.67 \ 0.67 \ 0.33)^t$$

and  $Q = (0.71 \ 0.71 \ 0)^t$

then the cosine similarity between D and Q  
will be :

$$\text{Sim}(D, Q) = \frac{0.67 \times 0.71 + 0.67 \times 0.71 + 0.33 \times 0}{\sqrt{(0.67^2 + 0.67^2 + 0.33^2)} \times \sqrt{(0.71^2 + 0.71^2 + 0^2)}}$$

## Summary: What's the point of using vector spaces?

---

- A well-formed algebraic space for retrieval
- Key: A user's query can be viewed as a (very) short document.
- Query becomes a vector in the same space as the docs.
- Can measure each doc's proximity to it.
- Natural measure of scores/ranking – no longer Boolean.
  - Queries are expressed as bags of words

## Interaction: vectors and phrases

---

- Scoring phrases doesn't fit naturally into the vector space world:
  - “**tangerine trees**” “**marmalade skies**”
  - Positional indexes don't calculate or store tf.idf information for “**tangerine trees**”
- Biword indexes (extend the idea to vector space)
  - For these, we can pre-compute tf.idf.
- We can use a positional index to boost or ensure phrase occurrence

## Vectors and wild cards

---

- How about the query ***tan\* marm\****?
  - Can we view this as a bag of words?
  - Thought: expand each wild-card into the matching set of dictionary terms.
- Danger – unlike the Boolean case, we now have *tfs* and *idfs* to deal with.



## Vector spaces and other operators

---

- Vector space queries are apt for no-syntax, bag-of-words queries (i.e. free text)
  - Clean metaphor for similar-document queries
- Not a good combination with Boolean, wild-card, positional query operators

# Query language vs. scoring

- May allow user a certain query language, say

---

  - Free text basic queries
  - Phrase, wildcard etc. in Advanced Queries.
- For scoring (oblivious to user) may use all of the above, e.g. for a free text query
  - Highest-ranked hits have query as a phrase
  - Next, docs that have all query terms near each other
  - Then, docs that have some query terms, or all of them spread out, with  $tf \times idf$  weights for scoring

## Exercises

---

- How would you augment the inverted index built in lectures 1–3 to support cosine ranking computations?
- Walk through the steps of serving a query.
- *The math of the vector space model is quite straightforward, but being able to do cosine ranking efficiently at runtime is nontrivial*

## Efficient cosine ranking

- Find the  $k$  docs in the corpus “nearest” to the query  $\Rightarrow k$  largest query-doc cosines.
- Efficient ranking:
  - Computing a single cosine efficiently.
  - Choosing the  $k$  largest cosine values efficiently.
    - Can we do this without computing all  $n$  cosines?
      - $n$  = number of documents in collection

## Limiting the accumulators: Frequency/impact ordered postings

---

- Idea: we only want to have accumulators for documents for which  $wf_{t,d}$  is high enough
- We sort postings lists by this quantity
- We retrieve terms by idf, and then retrieve only one block of the postings list for each term
- We continue to process more blocks of postings until we have enough accumulators
  - Can continue one that ended with highest  $wf_{t,d}$
  - The number of accumulators is bounded
- Anh et al. 2001

## Cluster pruning: preprocessing

---

- Pick  $\sqrt{n}$  *docs* at random: call these *leaders*
- For each other doc, pre-compute nearest leader
  - Docs attached to a leader: its *followers*;
  - Likely: each leader has  $\sim \sqrt{n}$  followers.


## Cluster pruning: query processing

---

- Process a query as follows:
  - Given query  $Q$ , find its nearest *leader*  $L$ .
  - Seek  $k$  nearest docs from among  $L$ 's followers.

# Visualization

---





# Why use random sampling

---

- Fast
- Leaders reflect data distribution

## General variants

---

- Have each follower attached to  $a=3$  (say) nearest leaders.
- From query, find  $b=4$  (say) nearest leaders and their followers.
- Can recur on leader/follower construction.

# Dimensionality reduction

---

- What if we could take our vectors and “pack” them into fewer dimensions (say 50,000→100) while preserving distances?
- (Well, almost.)
  - Speeds up cosine computations.
- Two methods:
  - Random projection.
  - “Latent semantic indexing”.


# Random projection onto $k \ll m$ axes

---

- Choose a random direction  $x_1$  in the vector space.
- For  $i = 2$  to  $k$ ,
  - Choose a random direction  $x_i$  that is orthogonal to  $x_1, x_2, \dots, x_{i-1}$ .
- Project each document vector into the subspace spanned by  $\{x_1, x_2, \dots, x_k\}$ .

# E.g., from 3 to 2 dimensions

---



$x_1$  is a random direction in  $(t_1, t_2, t_3)$  space.  
 $x_2$  is chosen randomly but orthogonal to  $x_1$ .

Dot product of  $x_1$  and  $x_2$  is zero.

## Guarantee

---

- With high probability, relative distances are (approximately) preserved by projection.
- Pointer to precise theorem in Resources.

# Computing the random projection

---

- Projecting  $n$  vectors from  $m$  dimensions down to  $k$  dimensions:
  - Start with  $m \times n$  matrix of terms  $\times$  docs,  $A$ .
  - Find random  $k \times m$  orthogonal projection matrix  $R$ .
  - Compute matrix product  $W = R \times A$ .
- $j^{\text{th}}$  column of  $W$  is the vector corresponding to doc  $j$ , but now in  $k \ll m$  dimensions.

## Cost of computation

---

- This takes a total of  $kmn$  multiplications.
- Expensive





## Latent semantic indexing (LSI)

---

- Another technique for dimension reduction
- Random projection was data-*independent*
- LSI on the other hand is data-*dependent*
  - Eliminate redundant axes
  - Pull together “related” axes – hopefully
    - ***car*** and ***automobile***